

EXPERIMENT 1

Shravanya Andhale
D15A-30

AIM:

To apply Decision Tree and Random Forest classification algorithms to analyze determinants of electricity satisfaction among Indian households and evaluate how digital, infrastructural, and socio economic factors influence behavioural perception within the energy governance system.

THEORY:

This experiment models electricity satisfaction as a behavioural outcome of energy service quality, infrastructure reliability, and digital engagement. Electricity satisfaction reflects institutional trust, perceived service quality, and governance effectiveness. Machine learning classification techniques are used to identify patterns that distinguish households with high satisfaction from those with low or moderate satisfaction.

1. Dataset Source

Dataset: India Residential Energy Survey conducted by the Council on Energy, Environment and Water.

Source link:

<https://www.kaggle.com/code/raitest/india-residential-energy-survey-ires-2020-eda/input?select=CEEW+-+IRES+Data.csv>

2. Dataset Description

The dataset contains household level information on electricity access, digital behavior, infrastructure reliability, appliance ownership, and socio economic characteristics.

Target Variable
q326_satis_electricity

Original scale:

- 1 Very Dissatisfied
- 2 Dissatisfied
- 3 Neutral
- 4 Satisfied
- 5 Very Satisfied

For classification, the variable was converted into binary form:

High Satisfaction 1 if value is 4 or 5

Low or Moderate Satisfaction 0 if value is 1, 2, or 3

After cleaning and removing missing values, the dataset contained 8563 observations. The test set consisted of 1713 observations with the following class distribution:

Low or Moderate Satisfaction 301

High Satisfaction 1412

Predictor Variables:-

Digital Variables:

q232_mobile_smart_n

q314_a_online_pay_ever_yn

Infrastructure Variables:

q312_grid_elec_meter_yn

q302_grid_hrs_no

q308_grid_voltage_low_app

Socio economic Variables:

asset_index_1

q208_priminc_earner_edu

q202_resp_age

q213_no_members

3. Mathematical Formulation of the Algorithm

Binary Classification

Let Y belong to {0,1} where

Y = 1 represents High Satisfaction

Y = 0 represents Low or Moderate Satisfaction

Decision Tree

Decision Tree partitions the feature space recursively based on impurity reduction. The Gini index is used as the splitting criterion:

Gini = 1 minus sum of squared class probabilities

$$\text{Gini} = 1 - \sum (p_i)^2$$

The tree selects splits that minimize Gini impurity at each node.

Random Forest

Random Forest is an ensemble learning method consisting of multiple decision trees built on bootstrap samples of the training data. At each split, a random subset of features is considered.

Final prediction is determined by majority voting across trees:

\hat{Y} equals mode of predictions from all trees

ROC AUC

Receiver Operating Characteristic curve measures model discrimination. AUC is computed as the area under the TPR versus FPR curve.

4. Algorithm Limitations

Decision Tree limitations

Sensitive to small data variations

Prone to overfitting

Limited ability to capture complex nonlinear relationships

Random Forest limitations

Reduced interpretability compared to single tree

Computationally intensive

Performance decreases when important predictors are missing

May still struggle with imbalanced datasets

Both algorithms

Cannot establish causality

Depend on quality and completeness of input variables

May not perform well when behavioral outcomes depend on unobserved psychological or contextual factors

5. Methodology Workflow

1. Load dataset
2. Inspect satisfaction distribution

3. Convert satisfaction into binary target variable
4. Select digital, infrastructure, and socio economic predictors
5. Remove missing values
6. Encode categorical variables
7. Perform stratified train test split
8. Train Decision Tree classifier
9. Train Random Forest classifier
10. Evaluate using Accuracy, Confusion Matrix, Classification Report, and ROC AUC
11. Apply class balancing and hyperparameter tuning for Random Forest
12. Performance Analysis

Decision Tree Results:

Accuracy 0.8179

AUC 0.5983

Confusion Matrix

```
[[ 9 292 ]  
 [ 20 1392 ]]
```

Classification Report

Class 0 Precision 0.31 Recall 0.03 F1 Score 0.05

Class 1 Precision 0.83 Recall 0.99 F1 Score 0.90

Interpretation

The Decision Tree model correctly classified most high satisfaction households but performed poorly in identifying low or moderate satisfaction households. The low recall of 0.03 for class 0 indicates severe difficulty in detecting dissatisfied households. AUC of 0.598 indicates weak discrimination capability.

Random Forest Results Before Tuning

Accuracy 0.8266

AUC 0.6702

Confusion Matrix

```
[[ 42 259 ]  
 [ 38 1374 ]]
```

Classification Report

Class 0 Precision 0.53 Recall 0.14 F1 Score 0.22

Class 1 Precision 0.84 Recall 0.97 F1 Score 0.90

Interpretation

Random Forest improved discrimination compared to Decision Tree. AUC increased to 0.67 indicating moderate ability to distinguish between satisfaction levels. Minority class detection improved compared to Decision Tree but remained limited.

Random Forest Results After Hyperparameter Tuning

Tuned RF Accuracy 0.8190

Tuned RF AUC 0.6788

Confusion Matrix

```
[[ 66 235 ]  
 [ 75 1337 ]]
```

Classification Report

Class 0 Precision 0.47 Recall 0.22 F1 Score 0.30

Class 1 Precision 0.85 Recall 0.95 F1 Score 0.90

Interpretation

After tuning and applying class balancing, recall for low or moderate satisfaction improved from 0.14 to 0.22. AUC increased from 0.6702 to 0.6788, indicating better overall discrimination. Although overall accuracy slightly decreased, minority class detection improved, demonstrating a better trade off between classes.

7. Hyperparameter Tuning

Hyperparameter tuning was performed using GridSearchCV with cross validation and ROC AUC as the scoring metric.

Parameters tuned included:

- Number of estimators
- Maximum depth
- Minimum samples split
- Minimum samples leaf
- Class weight set to balanced

The tuning process improved AUC and enhanced recall for the minority class, leading to a more balanced and robust classifier.

OUTPUT:

The Random Forest model after tuning achieved

Accuracy 0.8190

AUC 0.6788

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report, roc_curve, roc_auc_score

df = pd.read_csv("CEEW - IRES Data.csv", low_memory=False)
print("Dataset shape:", df.shape)

Dataset shape: (14851, 517)

df['q326_satis_electricity'].value_counts(dropna=False)

      count
q326_satis_electricity
    5.0      7648
    4.0      3614
    3.0      1145
    2.0      1055
    1.0       945
    NaN       444

dtype: int64
```

```
df_clean = df[df['q326_satis_electricity'].isin([1,2,3,4,5])].copy()

df_clean['High_Satisfaction'] = df_clean['q326_satis_electricity'].apply(
    lambda x: 1 if x >= 4 else 0
)

print("Target distribution:")
print(df_clean['High_Satisfaction'].value_counts())

*** Target distribution:
High_Satisfaction
1    11262
0     3145
Name: count, dtype: int64

X = df_clean[['q232_mobile_smart_n',           # smartphones
              'q314_a_online_pay_ever_yn',      # digital payment
              'q312_grid_elec_meter_yn',        # meter presence
              'q302_grid_hrs_no',             # power availability
              'q308_grid_voltage_low_app',    # voltage issue
              'asset_index_1',                # wealth
              'q208_priminc_earner_edu',      # education
              'q202_resp_age',               # age
              'q213_no_members']]            # household size

y = df_clean['High_Satisfaction']
```

```
data = pd.concat([X, y], axis=1).dropna()

X = data.drop('High_Satisfaction', axis=1)
y = data['High_Satisfaction']

print("Final class distribution:")
print(y.value_counts())

Final class distribution:
High_Satisfaction
1    7058
0     1505
Name: count, dtype: int64

X = pd.get_dummies(X, drop_first=True)

print("Feature shape after encoding:", X.shape)

Feature shape after encoding: (8563, 9)
```

```
X_train, X_test, y_train, y_test = train_test_split(  
    X, y, test_size=0.2, random_state=42, stratify=y)
```

```
print("Train distribution:")  
print(y_train.value_counts())
```

```
print("Test distribution:")  
print(y_test.value_counts())
```

```
Train distribution:  
High_Satisfaction  
1    5646  
0    1204  
Name: count, dtype: int64  
Test distribution:  
High_Satisfaction  
1    1412  
0     301  
Name: count, dtype: int64
```

```
dt = DecisionTreeClassifier(max_depth=5, random_state=42)  
dt.fit(X_train, y_train)
```

```
y_pred_dt = dt.predict(X_test)
```

```
print("Decision Tree Accuracy:", accuracy_score(y_test, y_pred_dt))
```

```
print("\nConfusion Matrix:")  
print(confusion_matrix(y_test, y_pred_dt))
```

```
print("\nClassification Report:")  
print(classification_report(y_test, y_pred_dt))
```

```
Decision Tree Accuracy: 0.8178633975481612
```

```
Confusion Matrix:  
[[ 9 292]  
 [ 20 1392]]
```

```
Classification Report:  
precision    recall   f1-score   support  
  
      0       0.31      0.03      0.05      301  
      1       0.83      0.99      0.90     1412  
  
accuracy          0.57      0.51      0.48     1713  
macro avg       0.74      0.82      0.75     1713  
weighted avg     0.74      0.82      0.75     1713
```

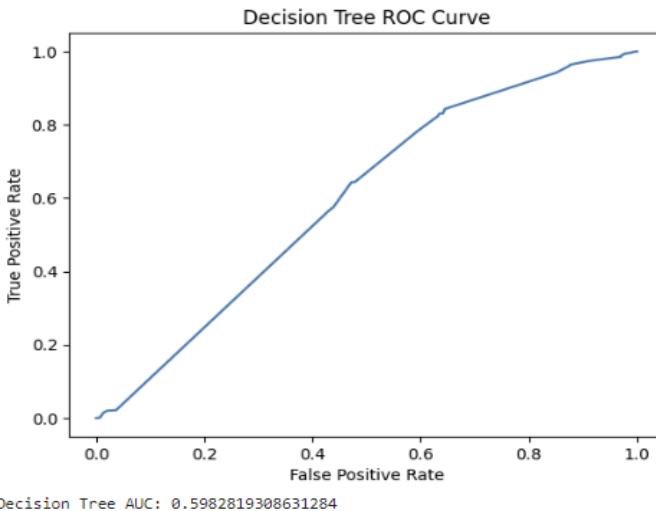
```

y_prob_dt = dt.predict_proba(X_test)[:,1]
fpr_dt, tpr_dt, _ = roc_curve(y_test, y_prob_dt)

plt.figure()
plt.plot(fpr_dt, tpr_dt)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Decision Tree ROC Curve")
plt.show()

print("Decision Tree AUC:", roc_auc_score(y_test, y_prob_dt))

```



Decision Tree AUC: 0.5982819308631284

```

rf_balanced = RandomForestClassifier(
    n_estimators=300,
    max_depth=None,
    min_samples_split=5,
    min_samples_leaf=2,
    class_weight='balanced',
    random_state=42
)

rf_balanced.fit(X_train, y_train)
y_pred_rf = rf_balanced.predict(X_test)

print("Random Forest Accuracy:", accuracy_score(y_test, y_pred_rf))
print(confusion_matrix(y_test, y_pred_rf))
print(classification_report(y_test, y_pred_rf))

y_prob_rf = rf_balanced.predict_proba(X_test)[:,1]
print("Random Forest AUC:", roc_auc_score(y_test, y_prob_rf))

```

	precision	recall	f1-score	support
0	0.46	0.21	0.28	301
1	0.85	0.95	0.90	1412
accuracy			0.82	1713
macro avg	0.65	0.58	0.59	1713
weighted avg	0.78	0.82	0.79	1713


Random Forest AUC: 0.6791137661995426

```
y_prob_rf = rf.predict_proba(X_test)[:,1]

fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)

plt.figure()
plt.plot(fpr_rf, tpr_rf)
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Random Forest ROC Curve")
plt.show()

print("Random Forest AUC:", roc_auc_score(y_test, y_prob_rf))
```



```

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [200, 400],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}

grid = GridSearchCV(
    RandomForestClassifier(class_weight='balanced', random_state=42),
    param_grid,
    cv=3,
    scoring='roc_auc',
    n_jobs=-1
)

grid.fit(X_train, y_train)

print("Best Parameters:", grid.best_params_)

Best Parameters: {'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 5, 'n_estimators': 400}

```

```

best_rf = grid.best_estimator_

y_pred_best = best_rf.predict(X_test)
y_prob_best = best_rf.predict_proba(X_test)[:,1]

print("Tuned RF Accuracy:", accuracy_score(y_test, y_pred_best))
print(confusion_matrix(y_test, y_pred_best))
print(classification_report(y_test, y_pred_best))
print("Tuned RF AUC:", roc_auc_score(y_test, y_prob_best))

Tuned RF Accuracy: 0.8190309398715704
[[ 66 235]
 [ 75 1337]]
              precision    recall   f1-score   support
      0       0.47     0.22     0.30      301
      1       0.85     0.95     0.90     1412

      accuracy                           0.82      1713
     macro avg       0.66     0.58     0.60      1713
  weighted avg       0.78     0.82     0.79      1713

Tuned RF AUC: 0.6788408327294289

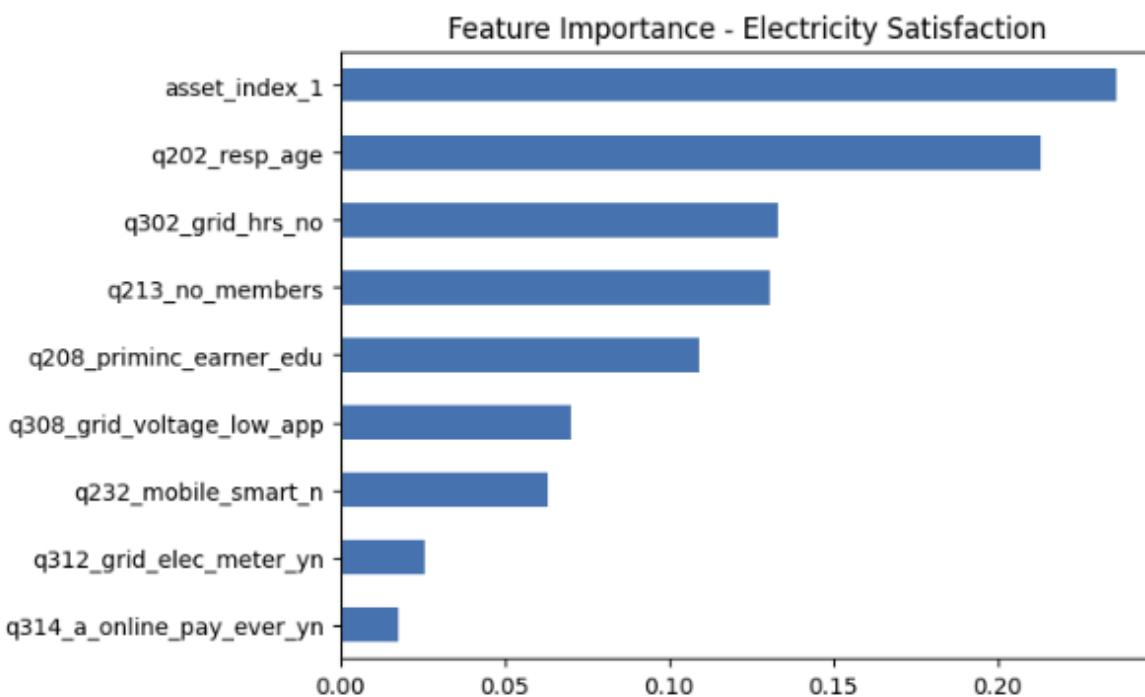
```

```

importances = rf.feature_importances_
feature_names = X.columns

feat_imp = pd.Series(importances, index=feature_names)
feat_imp.sort_values().plot(kind='barh')
plt.title("Feature Importance - Electricity Satisfaction")
plt.show()

```



CONCLUSION:

Decision Tree and Random Forest models were successfully applied to classify electricity satisfaction levels among Indian households. Random Forest outperformed Decision Tree in terms of AUC and minority class detection. The tuned Random Forest achieved an AUC of 0.6788 and improved recall for low satisfaction households.

The results indicate that electricity satisfaction is partially explained by digital participation, infrastructure reliability, and socio economic factors. However, moderate AUC values suggest that behavioral perception is influenced by additional unobserved factors.

This study demonstrates how machine learning techniques can be used to analyze behavioural sustainability and institutional trust within the electricity governance system, contributing to the understanding of digital transformation and energy service perception in Indian households.