

## **EXPERIMENT 4**

**Shravanya Andhale**  
**D15A-30**

**AIM:** Implement K-Nearest Neighbors (KNN) and evaluate model performance.

### **THEORY:**

#### **1. Dataset Source**

- Dataset: CEEW – India Residential Energy Survey (IRES) Microdata
- Source: Council on Energy, Environment and Water (CEEW)

Source link:

<https://www.kaggle.com/code/raitest/india-residential-energy-survey-ires-2020-eda/input?select=CEEW+-+IRES+Data.csv>

The dataset provides nationally representative household-level data on energy access, appliance ownership, awareness, billing, and socio-economic indicators.

#### **2. Dataset Description**

The dataset includes variables related to:

- Household asset index
- Education of primary income earner
- Respondent age
- Household size
- BEE star label awareness
- LED bulb ownership
- AC BEE star rating
- Refrigerator BEE star rating

Feature engineering was performed to create:

- **bee\_awareness** (1 if aware of BEE label)
- **led\_usage** (1 if LED bulbs present)
- **high\_star\_ac** (1 if AC rating  $\geq 4$ )
- **high\_star\_fridge** (1 if fridge rating  $\geq 4$ )

These were combined into:

- **Efficiency\_Score = sum of efficiency-positive behaviours**

Binary Target Variable:

- **Efficiency\_Binary = 1** if Efficiency\_Score  $\geq 3$  (High Efficiency)
- **Efficiency\_Binary = 0** otherwise (Low/Moderate Efficiency)

Test distribution:

- Class 0 (Low/Moderate): 2708
- Class 1 (High Efficiency): 263

This reflects real-world imbalance where strong sustainable adopters are fewer.

### 3. Mathematical Formulation of KNN

KNN is a non-parametric, distance-based classification algorithm.

- Step 1: Compute Euclidean distance between test point  $x$  and training point  $x_i$   
 $d(x, x_i) = \sqrt{\sum (x_j - x_{ij})^2}$
- Step 2: Select  $K$  nearest neighbors.
- Step 3: Assign class based on majority vote among neighbors.

Distance-weighted KNN assigns weight:

$$w_i = 1 / d(x, x_i)$$

Probability estimation:

$$P(y=1 | x) = (\text{Number of class 1 neighbors}) / K$$

Classification rule:

Predict 1 if  $P \geq$  threshold

Otherwise predict 0


### 4. Algorithm Limitations

- Sensitive to class imbalance
- Sensitive to feature scaling
- Struggles when classes overlap in feature space
- Computationally expensive for large datasets
- Cannot model complex nonlinear feature interactions

In this dataset, overlapping socio-economic characteristics reduce clean class separability.

## 5. Methodology / Workflow

- Import required libraries
- Load dataset
- Perform feature engineering to construct efficiency indicators
- Compute Efficiency\_Score
- Create binary target variable
- Select socio-economic predictor variables
- Handle missing values
- Perform stratified train-test split
- Apply feature scaling using StandardScaler
- Determine optimal K using accuracy vs K plot
- Train initial KNN model
- Evaluate using accuracy, confusion matrix, classification report, ROC-AUC
- Refine model using:
  - Smaller K
  - Distance weighting
  - Threshold adjustment (0.35)
- Re-evaluate using:
  - Balanced accuracy
  - Sensitivity and specificity
  - Precision-Recall curve
  - Threshold sensitivity analysis



## 6. Performance Analysis

### Initial Model (K = 19, threshold = 0.5)

- Accuracy: 0.9165
- ROC-AUC: 0.8459

#### Confusion Matrix:

[2686 22]
[ 226 37]]

Observations:

- Very high overall accuracy due to class imbalance.
- Recall for high-efficiency households was low (0.14).
- Strong ranking capability indicated by high AUC.

### Refined Model (K = 7, distance weighting, threshold = 0.35)

- Accuracy: 0.8792
- ROC-AUC: 0.7809

#### Confusion Matrix:

[2511 197]
[ 162 101]]

Metrics for High Efficiency (Class 1):

- Precision: 0.34
- Recall: 0.38
- F1-score: 0.36

Balanced accuracy improved.

Sensitivity significantly increased from **0.14 to 0.38**.

Although accuracy decreased slightly, minority detection improved meaningfully.

Precision-Recall analysis showed improved trade-off for minority class detection.

Threshold sensitivity analysis demonstrated controllable trade-off between recall and precision.

Interpretation:

- High-efficiency households are partially structurally distinguishable.
- However, significant overlap exists in socio-economic space.
- Behaviour is not fully determined by structural similarity.

## 7. Hyperparameter Tuning

The following parameters were tuned:

- **Number of neighbors (K)**
  - Tested K from 1 to 30
  - Initial optimal K: 19
  - Refined to K = 7 for better minority detection
- **Weights parameter**
  - Compared uniform vs distance weighting
  - Distance weighting improved minority sensitivity
- **Classification threshold**
  - Adjusted from 0.5 to 0.35
  - Increased recall for high-efficiency households

Hyperparameter tuning demonstrated that threshold calibration is essential for imbalanced behavioural classification problems.

## OUTPUT

- The KNN model achieved strong discrimination ability (AUC up to 0.8459).
- Refined model improved detection of high-efficiency households.
- Socio-economic predictors partially explain sustainable appliance adoption.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score,
    confusion_matrix,
    classification_report,
    roc_auc_score,
    roc_curve
)
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import (
    accuracy_score,
    confusion_matrix,
    classification_report,
    roc_auc_score,
    roc_curve
)
```

```

df_eff = df.copy()

df_eff['bee_awareness'] = np.where(
    df_eff['q401_bee_star_label_heard_yn'] == 1, 1, 0
)

df_eff['led_usage'] = np.where(
    df_eff['q405_c_led_bulb_no'] > 0, 1, 0
)

df_eff['high_star_ac'] = np.where(
    df_eff['q434_ac_most_bee_rating'] >= 4, 1, 0
)

df_eff['high_star_fridge'] = np.where(
    df_eff['q469_fridge_most_bee_star'] >= 4, 1, 0
)

df_eff['Efficiency_Score'] = (
    df_eff['bee_awareness'] +
    df_eff['led_usage'] +
    df_eff['high_star_ac'] +
    df_eff['high_star_fridge']
)

df_eff['Efficiency_Score'].value_counts()

```

	count
Efficiency_Score	
1	7878
2	3693
0	1963
3	1200
4	117

```

df_eff['Efficiency_Binary'] = np.where(
    df_eff['Efficiency_Score'] >= 3, 1, 0
)

print("Target distribution:")
print(df_eff['Efficiency_Binary'].value_counts())

```

Target distribution:

Efficiency_Binary	count
0	13534
1	1317

Name: count, dtype: int64

```

predictors = [
    'asset_index_1',
    'q208_priminc_earner_edu',
    'q202_resp_age',
    'q213_no_members'
]

df_model = df_eff[predictors + ['Efficiency_Binary']].dropna()

X = df_model[predictors]
y = df_model['Efficiency_Binary']

print("Final dataset shape:", df_model.shape)

```

Final dataset shape: (14851, 5)

```

X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print("Train distribution:")
print(y_train.value_counts())

print("Test distribution:")
print(y_test.value_counts())

Train distribution:
Efficiency_Binary
0    10826
1     1054
Name: count, dtype: int64
Test distribution:
Efficiency_Binary
0    2708
1     263
Name: count, dtype: int64

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

```


```

for k in range(1, 31):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    y_pred = knn.predict(X_test_scaled)
    accuracies.append(accuracy_score(y_test, y_pred))

plt.figure(figsize=(8,5))
plt.plot(range(1,31), accuracies)
plt.xlabel("K value")
plt.ylabel("Accuracy")
plt.title("KNN Accuracy for Different K")
plt.show()

optimal_k = np.argmax(accuracies) + 1
print("Optimal K:", optimal_k)

```



## Initial KNN Results:

```
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_train_scaled, y_train)

y_pred = knn.predict(X_test_scaled)
y_prob = knn.predict_proba(X_test_scaled)[:,1]

print("KNN Accuracy:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
print(cm)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("ROC-AUC:", roc_auc_score(y_test, y_prob))

KNN Accuracy: 0.9165264220801077


Confusion Matrix:
[[2686  22]
 [ 226  37]]


Classification Report:
              precision    recall   f1-score   support
          0       0.92      0.99      0.96     2708
          1       0.63      0.14      0.23      263

      accuracy           0.92      2971
      macro avg       0.77      0.57      0.59      2971
  weighted avg       0.90      0.92      0.89      2971

ROC-AUC: 0.8459619715699436
```

```
plt.figure(figsize=(6,5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix Heatmap")
plt.show()
```





## Reduced k=7, added distance as weights

```

knn = KNeighborsClassifier(
    n_neighbors=7,      # smaller K
    weights='distance' # distance weighting
)

knn.fit(X_train_scaled, y_train)

y_prob = knn.predict_proba(X_test_scaled)[:,1]

# threshold = 0.35  # try 0.3-0.4
y_pred = (y_prob >= threshold).astype(int)

print("Accuracy:", accuracy_score(y_test, y_pred))

cm = confusion_matrix(y_test, y_pred)
print("\nConfusion Matrix:")
print(cm)

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("ROC-AUC:", roc_auc_score(y_test, y_prob))

```

Accuracy: 0.8791652642208011

Confusion Matrix:

[2511 197]
[ 162 101]

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.93	0.93	2708
1	0.34	0.38	0.36	263
accuracy			0.88	2971
macro avg	0.64	0.66	0.65	2971
weighted avg	0.89	0.88	0.88	2971

ROC-AUC: 0.7808823033849852

## Balanced accuracy:

```
from sklearn.metrics import balanced_accuracy_score

# Balanced Accuracy
bal_acc = balanced_accuracy_score(y_test, y_pred)
print("Balanced Accuracy:", bal_acc)

# Sensitivity (Recall for class 1)
sensitivity = cm[1,1] / (cm[1,1] + cm[1,0])
print("Sensitivity (Recall for High Efficiency):", sensitivity)

# Specificity (Recall for class 0)
specificity = cm[0,0] / (cm[0,0] + cm[0,1])
print("Specificity (Low/Moderate Detection):", specificity)


Balanced Accuracy: 0.6556415015922404
Sensitivity (Recall for High Efficiency): 0.3840304182509506
Specificity (Low/Moderate Detection): 0.9272525849335302
```

```
from sklearn.metrics import precision_recall_curve, average_precision_score

precision, recall, thresholds = precision_recall_curve(y_test, y_prob)

plt.figure(figsize=(6,5))
plt.plot(recall, precision)
plt.xlabel("Recall")
plt.ylabel("Precision")
plt.title("Precision-Recall Curve")
plt.show()

ap_score = average_precision_score(y_test, y_prob)
print("Average Precision Score:", ap_score)
```




## Impact of threshold:

```
thresholds_test = np.linspace(0.1, 0.6, 30)

recall_list = []
precision_list = []

for t in thresholds_test:
    y_temp = (y_prob >= t).astype(int)
    recall_list.append(
        classification_report(y_test, y_temp, output_dict=True)['1']['recall']
    )
    precision_list.append(
        classification_report(y_test, y_temp, output_dict=True)['1']['precision']
    )

plt.figure(figsize=(7,5))
plt.plot(thresholds_test, recall_list, label="Recall (High Efficiency)")
plt.plot(thresholds_test, precision_list, label="Precision (High Efficiency)")
plt.xlabel("Threshold")
plt.ylabel("Score")
plt.legend()
plt.title("Threshold Sensitivity Analysis")
plt.show()
```



## CONCLUSION

- High energy-efficiency adoption exhibits partial structural clustering.
- However, sustainable behaviour is not fully separable in socio-economic feature space.
- Structural factors alone cannot completely explain energy-efficient behaviour.
- Threshold calibration improves practical and policy relevance.

This experiment demonstrates that sustainable household behaviour is partially predictable but influenced by additional contextual, informational, and behavioural factors beyond structural similarity.